

Detecting and Characterizing Events

Anonymous EMNLP submission

Abstract

Significant events are characterized by interactions between entities (e.g., countries, organizations, individuals) that deviate from typical interaction patterns. Investigators, such as historians, commonly read large quantities of text to construct an accurate picture of who, what, when, and where and event happened. In this work, we present the *Capsule* model for analyzing documents to identify and characterize events of potential significance. Specifically, we develop a model based on topic modeling to distinguish between topics that describe “business-as-usual” and topics that deviate from these patterns. To demonstrate this model, we analyze real-world datasets, including a corpus of over 2 million US State Department cables from the 1970s; we provide open-source implementations of an inference algorithm for the Capsule model and a pipeline to visualize and explore its results.

1 Introduction

Events are difficult to define; historians and political scientists read large quantities of text to construct an accurate picture of a single historical event. Events are interesting by definition: they are the hidden causes of anomalous observations. But they are also inherently abstract—we can observe that changes occur, but we cannot directly observe whether or not an event occurs.

Consider embassies sending diplomatic messages, such as shown in Figure 1. The Bangkok embassy, Hong Kong Embassy, and the State Department all have *typical concerns* about which they usually send

messages. At date d , however, the message content changes for all three entities—again, we only observe the changes in message content, and do not observe the event directly. Our first goal is to determine *when* events happen, or identify these rare but pervasive deviations from the typical concerns.

Our second goal is to characterize *what* occurs. We rely on topic modeling (Blei, 2012) to summarize content and to characterize events.

We develop a Bayesian model that discovers the typical concerns of authors, identifies when events occur, and characterizes these events; we call this the *Capsule* model, as it encapsulates events.

We first review previous research related to event detect, summarization, and visualization. In Section 2, we describe the Capsule model and how to infer the latent parameters (the appendix provides further inference details). Section 3 provides an exploration of results on a real-world dataset.

Related work. We first review previous work on automatic event detection and other related concepts.

In both univariate and multivariate settings, the goal is often the same: analysts want to predict whether or not a rare events will occur (Weiss and Hirsh, 1998; Das et al., 2008). Capsule, in contrast, is designed to help analysts explore and understand the original data: our goal is interpretability, not prediction.

A common goal is to identify clusters of documents; these approaches are used on news articles (Zhao et al., 2012; Zhao et al., 2007; Zhang et al., 2002; Li et al., 2005; Wang et al., 2007; Allan et al., 1998) and social media posts (VanDam, 2012; Lau et al., 2012; Jackoway et al., 2011; Sakaki et al., 2010; Reuter and Cimiano, 2012; Becker et al., 2010; Sayyadi et al., 2009). In the case of news articles, the task is to create new clusters as novel news stories appear—this does not help disentangle typical content from rare events of interest. Social media approaches identify rare events, but the methods are designed for short, noisy documents; they are not appropriate for larger documents that contain information about a variety of subjects.

Many existing methods use document terms as features, frequently weighted by tf-idf value (Fung et al., 2005; Kumaran and Allan, 2004; Brants et al., 2003; Das Sarma et al., 2011; Zhao et al., 2007; Zhao et al., 2012); here, events are bursts in groups of terms.

Topic models (Blei, 2012) reduce the dimensionality of text data; they have been used to help detect events mentioned in social media posts (Lau et al., 2012; Dou et al., 2012) and posts relevant to monitored events (VanDam, 2012). We rely on topic models to characterize both typical content and events, but grouped observations can also be summarized directly (Peng et al., 2007; Chakrabarti and Punera, 2011; Gao et al., 2012).

In addition to text data over time, author (Zhao et al., 2007), news outlet (Wang et al., 2007), and spatial information (Neill et al., 2005; Mathioudakis et al., 2010; Liu et al., 2011) can be used to augment event detection. Capsule uses author information in order to characterize typical concerns of authors.

Detecting and characterizing relationships (Schein et al., 2015; Linderman and Adams, 2014; Das Sarma et al., 2011) is related to event detection. When a message recipient is known, Capsule’s author input can be replaced with a sender-receiver pair, but the model could be further tailored for interactions within networks.

2 The Capsule Model

In this section we develop the Capsule model. Capsule captures patterns in entity behavior and identifies

time intervals in which many entities deviate from these patterns. The model relies on high dimensional entity behavior data over time, such as text messages being sent between entities. We first review topic models at a high level, and give the intuition on Capsule. Then, we formally specify our model and describe how we learn the hidden variables.

Background: Topic Models. Capsule relies on topic models to model text data. Topic models are algorithms for discovering the main themes in a large collection of documents; each document can then be summarized in terms of the global themes. More formally, a topic k is a probability distribution over the set of vocabulary words. Each document d is represented as a distribution over topics θ_d . Thus we can imagine that when we generate a document, we first pick which topics are relevant (and in what proportions). Under the LDA topic model (Blei et al., 2003), we know the number of words in each document. Then, for each word, we select a single topic from this distribution over topics, and finally select a vocabulary term from the corresponding topic’s distribution over the vocabulary. Alternatively, we can cast topic modeling as factorization, such as in Poisson factorization (Gopalan et al., 2014), and draw a word count for each term in the vocabulary.

The Capsule Model. Topic models are often applied to provide a structure for an otherwise unstructured collection of documents. Documents, however, are often accompanied by metadata, such as the date written or author attribution; this information is not exploited by traditional topic models. The Capsule model uses both author and date information to identify and characterize events that influence the content of the collection.

Consider an entity like the Bangkok American embassy, shown in Figure 1. We can imagine that there is a stream of messages (or *diplomatic cables*) being sent by this embassy—some might be sent to the US State Department, others to another American embassy like Hong Kong. An entity will usually talk about certain topics; the Bangkok embassy, for instance, is concerned with topics regarding southeast Asia more generally.

Now imagine that at a particular time t , an event occurs, such as the capture of Saigon during the Vietnam war. We do not directly observe that events occur, but we do observe the message stream. Using




Figure 1: Cartoon intuition of Capsule. The Bangkok embassy, Honk Kong embassy, and State Department all have typical concerns about which they usually send messages. When an event occurs at time t , the stream of message content alters to include the event, then fades back to “business as usual.” Capsule discovers both entities’ typical concerns and the event locations and content.

this stream, each event be described as a distribution over the vocabulary, similar to how topics are distributions over these same terms. When an event occurs, the message content changes for multiple entities. The day following the capture of Saigon, the majority of the diplomatic cables sent by the Bangkok embassy were about Vietnam war refugees. Thus we imagine that an entity’s stream of messages is controlled by what it usually talks about as well as the higher level stream of unobserved events.

Model Specification. We formally describe Capsule. The observed data are word counts $w_{d,v}$ for document d and vocabulary term v ; each document d also has an author (or entity) a_d and a time (or date) interval i_d associated with it.

The hidden variables of this model are general topics of conversation β , authors’ typical concerns ϕ , event descriptions π , event strengths ψ , and document-specific topics θ and event relevancy ϵ . The graphical model in Figure 2 shows considers dependencies between these latent parameters and the observed data.

As in topic modeling, we represent the topics of conversation with a $K \times V$ matrix β , where K is a low dimensional number of topics that we wish to capture, and V is the size of our vocabulary; each row β_k is normalized such that it represents




Figure 2: A directed graphical model of Capsule to show considered dependencies. Shaded nodes w are observed word counts. Unshaded nodes are hidden variables—at the global level, these are topics β , entity typical concerns ϕ , interval strength ϵ , and interval content descriptions π . At the document level, we have hidden variables for local entity concerns θ and interval relevancy ϵ . Plates denote replication: there are D documents, T intervals, N entities, K topics, and V vocabulary terms. Fixed hyperparameters are indicated by black squares.

the probability of seeing vocabulary word v when discussing topic k . As a generative process, we draw these general topics from a Dirichlet distribution, or $\beta_k \sim \text{Dirichlet}_V(\alpha_\beta)$.

The concerns of author n are represented with ϕ_n , a K -dimensional topic vector, where each element is drawn from a gamma distribution, or $\phi_{n,k} \sim \text{Gamma}(s_\phi, r_\phi)$.¹ Similar to topic modeling, we represent the contents of each document in topic space; each document d has a K -dimensional latent parameter θ_d to describe the particular contents of that document. Unlike traditional topic models, each document d ’s topics depend on the concerns of the author a_d ; each document topic $\theta_{d,k}$ is drawn from a gamma distribution parameterized by the corresponding author concerns $\phi_{a_d,k}$: $\theta_{d,k} \sim \text{Gamma}(s_\theta, \phi_{a_d,k})$.

To represent events, we consider discrete intervals of time. Each interval t has a corresponding interval strength ψ_t and description π_t . Event strengths are a single value for each interval t , and are drawn from a gamma distribution: $\psi_{n,k} \sim \text{Gamma}(s_\psi, r_\psi)$. These strengths indicate how important the interval is in determining message content. Interval descriptions are similar to topics: each description is a V -dimensional vector drawn from a Dirichlet distribution over the vocabulary terms, or $\pi_k \sim$

¹We use the shape-rate parameterization for all Gamma distributions.

Dirichlet $_V(\alpha_\pi)$.

Just as we describe each document d in terms of relevant topics with the θ_d parameters, we also describe the relevancy of each time interval with the ϵ_d parameters. These interval relevancy parameters are drawn from gamma distributions and depend on the overall strength ψ of the corresponding interval; for interval t and document d (written at time i_d), we have $\epsilon_{d,t} \sim \text{Gamma}(s_\epsilon, \psi_{i_d,t})$.

Conditional on the hidden variables and the author and time metadata, Capsule is a model of how document word counts came to be. For document d and vocabulary term v , we generate the word counts form a Poisson distribution parameterized by the documents topics θ_d and relevant events ϵ , as well as global topic β and event descriptions π :

$$w_{d,v} \sim \text{Poisson} \left(\theta_d^\top \beta_v + \sum_{t=1}^T f(i_d, t) \epsilon_{d,t} \pi_{t,v} \right), \quad (1)$$

where f is some function of decay. This function is important because events should not remain at their full strength indefinitely, but should decay over time. In our experiments, we consider step functions, linear decay, and exponential decay. Figure 3 gives the full generative process for Capsule.

Learning the hidden variables. In order to use the Capsule model to explore the observed documents, we must compute the posterior distribution. Conditional on the observed word counts w , our goals to to compute the posterior values of the hidden parameters—global interval strengths ψ , interval descriptions π , entity concerns ϕ , and topics β , as well as document-specific entity concerns θ and interval relevancy parameters ϵ .

As for many Bayesian models, the exact posterior for Capsule is not tractable to compute; approximating it is our central statistical and computational problem. We develop an approximate inference algorithm for Capsule based on variational methods (Wainwright and Jordan, 2008).²

Variational inference approaches the problem of posterior inference by minimizing the KL divergence from an approximating distribution q to the true pos-

- for each time step $t = 1:T$,
 - draw interval description over vocabulary $\pi_t \sim \text{Dirichlet}_V(\alpha)$
 - draw interval strength $\psi_t \sim \text{Gamma}(s_\psi, r_\psi)$
- for each topic $k = 1:K$,
 - draw topic distribution over vocabulary $\beta_k \sim \text{Dirichlet}_V(\alpha)$
 - for each entity $n = 1:N$,
 - draw entity concern $\phi_{n,k} \sim \text{Gamma}(s_\phi, r_\phi)$
- for each document $d = 1:D$ sent at time i_d by author a_d ,
 - for each topic $k = 1:K$,
 - draw local entity concern $\theta_{d,k} \sim \text{Gamma}(s_\theta, \phi_{a_d,k})$
 - for each time $t = 1:T$,
 - draw local interval relevancy $\epsilon_{d,t} \sim \text{Gamma}(s_\epsilon, \psi_{i_d,t})$
 - for each vocabulary term $v = 1:V$,
 - draw word counts $w_{d,v} \sim \text{Poisson} \left(\theta_d^\top \beta_v + \sum_{t=1}^T f(i_d, t) \epsilon_{d,t} \pi_{t,v} \right)$

Figure 3: The generative process for Capsule.

terior p . This is equivalent to maximizing the ELBO,

$$\mathcal{L}(q) = \mathbb{E}_{q(\psi, \pi, \phi, \beta, \theta, \epsilon)} [\log p(w, \psi, \pi, \phi, \beta, \theta, \epsilon) - \log q(\psi, \pi, \phi, \beta, \theta, \epsilon)]. \quad (2)$$

We define the approximating distribution q using the mean field assumption:

$$q(\psi, \pi, \phi, \beta, \theta, \epsilon) = \prod_{t=1}^T \left[q(\pi_t | \lambda_t^\pi) q(\psi_t | \lambda_t^\psi) \right] \prod_{k=1}^K \left[q(\beta_k | \lambda_k^\beta) \prod_{n=1}^N q(\phi_{n,k} | \lambda_{n,k}^\phi) \right] \prod_{d=1}^D \left[\prod_{k=1}^K q(\theta_{d,k} | \lambda_{d,k}^\theta) \prod_{t=1}^T q(\epsilon_{d,t} | \lambda_{d,t}^\epsilon) \right] \quad (3)$$

The variational distributions $q(\pi)$ and $q(\beta)$ are both Dirichlet-distributed with free variational parameters λ^π and λ^β , respectively. Similarly, $q(\psi)$, $q(\phi)$, $q(\theta)$ and $q(\epsilon)$ are all gamma-distributed with

²Source code will be released on github upon publication.

corresponding free variational parameters λ^ψ , λ^ϕ , λ^θ , and λ^ϵ .

The expectations under q , which are needed to maximize the ELBO, have closed form analytic updates, as detailed in Appendix A. We update each parameter in turn, following standard coordinate ascent variational inference techniques. Full details on our inference algorithm can be found in the appendix. This algorithm produces a fitted variational distribution which can then be used as a proxy for the true posterior, allowing us to explore a collection of documents with Capsule.

3 Evaluation

In this section we explore the performance of Capsule on a collection of US State Department cables. These cables were sent between 1973 and 1977 and obtained from the History Lab at Columbia,³ which received them from the Central Foreign Policy Files at the National Archives. In addition to the text of the cables themselves, each document is supplemented with information about who sent the cable (e.g., the State Department, the U.S. Embassy in Saigon, or an individual by name), who received the cable (often multiple entities), and the date the cable was sent. To test our model, we considered only cables from 1976. Using a vocabulary of size 6293, we omitted cables with fewer than three terms, resulting in a collection of 335,631 messages sent between 542 entities. We selected a weekly duration for the time intervals, as few cables were sent on the weekends.

We fit Capsule with 100 topics and using an exponential decay with mean lifetime of 3—this indicates that most intervals would no longer be relevant after about 3 weeks. For any given week, we can sort the documents by their interval relevancy parameters ϵ . Table 1 show the top documents for the week of July, 6, 1976. Two weeks prior, an Air France airplane had been hijacked and taken to Entebbe, Uganda. The flight had originally be scheduled to travel between Tel Aviv and Paris, and on July 4th, Israeli operatives recovered the passengers who were being held hostage. Our model accurately recovers cables that are relevant to this event, allowing investigators to use it as a tool to uncover relevant documents for a given time interval.

³<http://history-lab.org>




Figure 4: TODO




Figure 5: TODO

The week of December 12, 1976, the State Department sends a query for all posts to report on the presence of gambling equipment, or “gaming devices.” Nearly all posts respond that week with a report. Unsurprisingly, the top terms in the event description π for that week are the terms *gaming* and *devices*.

These exploratory results show that our model is successfully capturing when multiple entities are discussing the same subjects and that our model can be used to explore the underlying data by providing a structured scaffold from which to view the data.

A Inference

In this appendix, we describe the details of the variational inference algorithm for Capsule. This algorithm fits the parameters of the variational distribution q in Eq. 3 so that it is close in KL divergence to the posterior.

Recall that the variational distributions $q(\pi)$ and $q(\beta)$ are both Dirichlet-distributed with free variational parameters λ^π and λ^β , respectively. Similarly, the variational distributions $q(\psi)$, $q(\phi)$, $q(\theta)$ and $q(\epsilon)$ are all gamma-distributed with corresponding free variational parameters λ^ψ , λ^ϕ , λ^θ , and λ^ϵ . For these gamma-distributed variables, each free parameter λ has two components: shape s and rate r .

ϵ	date	sender	recipient	subject
0.04133056	1976-07-06	PEKING	STATE	POSSIBLE SC MEETING ON ISRAELI RESCUE
0.04095504	1976-07-09	MASERU	STATE	UGANDAN ROLE IN AIR FRANCE HIJACKING
0.03768492	1976-07-04	STATE	ABU DHABI	ISRAELI RESCUE OPERATION
0.03583023	1976-07-14	STATE	BERN	VIOLATIONS OF SEC REGULATIONS ALLEG...
0.02769182	1976-07-06	ROME	STATE	POSSIBLE SC MEETING ON ISRAELI RESCUE
0.02531757	1976-07-12	CARACAS	STATE	SECURITY COUNCIL DEBATE ON ENTEBBE...

Table 1: Top cables by ϵ for the week of July 6, 1976. The majority of these cables are concerning the Israeli rescue of a Air France airplane hijacking that had occurred the week prior.

data	full Capsule model	events only	topics only	entities only
US State Dept. Cables	-2.41e7	-1.82e7	-2.69e7	-2.43e7
arXiv	-3.12e6	-3.57e6	-3.14e6	-3.57e6

Table 2: Held-out data log likelihood.

Minimizing the KL divergence between the true posterior p and the variational approximation q is equivalent to maximizing the ELBO (Eq. 2). We achieve this with closed form coordinate updates, as the Capsule model is specified with the required conjugate relationships that make this approach possible (Ghahramani and Beal, 2001).

To obtain simple updates, we first rely on auxiliary latent variables z . These variables, when marginalized out, leave the original model intact. The Poisson distribution has an additive property; specifically if $w \sim \text{Poisson}(a + b)$, then $w = z_1 + z_2$, where $z_1 \sim \text{Poisson}(z_1)$ and $z_2 \sim \text{Poisson}(z_2)$. We apply this decomposition to the word count distribution in Eq. 1 and define Poisson variables for each component of the word count:

$$z_{d,v,k}^{\mathcal{K}} \sim \text{Poisson}(\theta_{d,k} \beta_{k,v})$$

$$z_{d,v,t}^{\mathcal{T}} \sim \text{Poisson}(f(i_d, t) \epsilon_{d,t} \pi_{t,v}).$$

The \mathcal{K} and \mathcal{T} superscripts indicate the contributions from entity concerns and events, respectively. Given these variables, the total word count is deterministic:

$$w_{d,v} = \sum_{k=1}^K z_{d,v,k}^{\mathcal{K}} + \sum_{t=1}^T z_{d,v,t}^{\mathcal{T}}.$$

Coordinate-ascent variational inference is derived from complete conditionals, i.e., the conditional distributions of each variable given the other variables

and observations. These conditionals define both the form of each variational factor and their updates. The following are the complete conditional for each of the gamma- and Dirichlet-distributed latent parameters. The notation $D(i)$ is used for the set of documents sent by entity i ; $D(t)$ is the set of documents sent impacted by events at time t (e.g., all documents after the event in the case of exponential decay).

$$\pi_t \mid \mathbf{W}, \psi, \phi, \beta, \theta, \epsilon, z \sim$$

$$\text{Dirichlet}_V \left(\alpha_\pi + \sum_{d=1}^D \langle z_{d,1,t}^{\mathcal{T}}, \dots, z_{d,V,t}^{\mathcal{T}} \rangle \right) \quad (4)$$

$$\beta_k \mid \mathbf{W}, \psi, \pi, \phi, \theta, \epsilon, z \sim$$

$$\text{Dirichlet}_V \left(\alpha_\beta + \sum_{d=1}^D \langle z_{d,1,k}^{\mathcal{K}}, \dots, z_{d,V,k}^{\mathcal{K}} \rangle \right) \quad (5)$$

$$\psi_t \mid \mathbf{W}, \pi, \phi, \beta, \theta, \epsilon, z \sim$$

$$\text{Gamma} \left(s_\psi + |D(t)| s_\epsilon, r_\psi + \sum_{d \in D(t)} \epsilon_{d,t} \right) \quad (6)$$

$$\phi_{i,k} \mid \mathbf{W}, \psi, \pi, \beta, \theta, \epsilon, z \sim$$

$$\text{Gamma} \left(s_\phi + |D(i)| s_\theta, r_\phi + \sum_{d \in D(i)} \theta_{d,k} \right) \quad (7)$$

$$\theta_{d,k} \mid \mathbf{W}, \psi, \pi, \phi, \beta, \epsilon, z \sim \text{Gamma} \left(s_\theta + \sum_{v=1}^V z_{d,v,k}^{\mathcal{K}}, \phi_{a_d,k} + \sum_{v=1}^V \beta_{k,v} \right) \quad (8)$$

$$\epsilon_{d,t} \mid \mathbf{W}, \psi, \pi, \phi, \beta, \theta, z \sim \text{Gamma} \left(s_\epsilon + \sum_{v=1}^V z_{d,v,t}^{\mathcal{T}}, \psi_t + f(i_d, t) \sum_{v=1}^V \pi_{t,v} \right) \quad (9)$$

The complete conditional for the auxiliary variables has the form $z_{d,v} \mid \psi, \pi, \phi, \beta, \theta, \epsilon \sim \text{Mult}(w_{d,v}, \omega_{d,v})$, where

$$\omega_{d,v} \propto \langle \theta_{d,1} \beta_{1,v}, \dots, \theta_{d,K} \beta_{K,v}, f(i_d, 1) \epsilon_{d,1} \pi_{1,v}, \dots, f(i_d, T) \epsilon_{d,T} \pi_{T,v} \rangle. \quad (10)$$

Intuitively, these variables allocate the data to one of the entity concerns or events, and thus can be used to explore the data.

Given these conditionals, the algorithm sets each parameter to the expected conditional parameter under the variational distribution. The mean field assumption guarantees that this expectation will not involve the parameter being updated. Algorithm 1 shows our variational inference algorithm.

References

- James Allan, Ron Papka, and Victor Lavrenko. 1998. Online new event detection and tracking. In *Proceedings of the 21st annual international ACM SIGIR conference on Research and development in information retrieval*, pages 37–45. ACM.
- Hila Becker, Mor Naaman, and Luis Gravano. 2010. Learning similarity metrics for event identification in social media. In *Proceedings of the third ACM international conference on Web search and data mining*, pages 291–300. ACM.
- David M. Blei, Andrew Y. Ng, and Michael I. Jordan. 2003. Latent Dirichlet allocation. *JMLR*, 3:993–1022, March.

⁴ $V(d)$ is the set of vocabulary indices for the collection of words in document d . We could also iterate over all V , but as zero word counts give $\mathbb{E}[z_{d,v}] = 0 \forall v \notin V(d)$, the two are equivalent.

Algorithm 1: Variational Inference for Capsule

Input: word counts w

Output: approximate posterior of latent parameters $(\psi, \pi, \phi, \beta, \theta, \epsilon)$ in terms of variational parameters $\lambda = \{\lambda^\psi, \lambda^\pi, \lambda^\phi, \lambda^\beta, \lambda^\theta, \lambda^\epsilon\}$

Initialize $\mathbb{E}[\beta]$ to slightly random around uniform

Initialize $\mathbb{E}[\psi], \mathbb{E}[\pi], \mathbb{E}[\phi], \mathbb{E}[\theta], \mathbb{E}[\epsilon]$ to uniform

for iteration $m = 1 : M$ **do**

set $\lambda^\psi, \lambda^\pi, \lambda^\phi, \lambda^\beta, \lambda^\theta, \lambda^\epsilon$ to respective priors, excluding $\lambda^{\theta, rate}$ and $\lambda^{\epsilon, rate}$, which are set to 0

update $\lambda^{\theta, rate} += \sum_V \mathbb{E}[\beta_v]$

for each document $d = 1 : D$ **do**

for each term $v \in V(d)$ ⁴ **do**

set $(K + T)$ -vector $\omega_{d,v}$ using $\mathbb{E}[\pi]$, $\mathbb{E}[\theta]$, and $\mathbb{E}[\epsilon]$, as shown in Eq. 10

set $(K + T)$ -vector $\mathbb{E}[z_{d,v}] = w_{d,v} * \omega_{d,v}$

update $\lambda_d^{\theta, shape} += \mathbb{E}[z_{d,v}^{\mathcal{K}}]$ (Eq. 8)

update $\lambda_d^{\epsilon, shape} += \mathbb{E}[z_{d,v}^{\mathcal{T}}]$ (Eq. 9)

update $\lambda_v^\beta += \mathbb{E}[z_{d,v}^{\mathcal{K}}]$ (Eq. 5)

update $\lambda_v^\pi += \mathbb{E}[z_{d,v}^{\mathcal{T}}]$ (Eq. 4)

end

update $\lambda_d^{\theta, rate} += \mathbb{E}[\phi_{a_d}]$ (Eq. 8)

update $\lambda_d^{\epsilon, rate} += \mathbb{E}[\psi]$ (Eq. 9)

set $\mathbb{E}[\theta_d] = \lambda_d^{\theta, shape} / \lambda_d^{\theta, rate}$

set $\mathbb{E}[\epsilon_d] = \lambda_d^{\epsilon, shape} / \lambda_d^{\epsilon, rate}$

update $\lambda_{a_d}^{\phi, shape} += s_\theta$ (Eq. 7)

update $\lambda_t^{\psi, shape} += s_\epsilon \forall t : f(i_d, t) \neq 0$ (Eq. 6)

update $\lambda_{a_d}^{\phi, rate} += \theta_d$ (Eq. 7)

update $\lambda^{\psi, rate} += \epsilon_d$ (Eq. 6)

end

set $\mathbb{E}[\phi] = \lambda^{\phi, shape} / \lambda^{\phi, rate}$

set $\mathbb{E}[\beta_k] = \lambda^{\beta, v} / \sum_v \lambda^{\beta_k} \forall k$

set $\mathbb{E}[\psi] = \lambda^{\psi, shape} / \lambda^{\psi, rate}$

set $\mathbb{E}[\pi_t] = \lambda^{\pi, v} / \sum_v \lambda^{\pi_t} \forall t$

end

return λ

- David M Blei. 2012. Probabilistic topic models. *Communications of the ACM*, 55(4):77–84.
- Thorsten Brants, Francine Chen, and Ayman Farahat. 2003. A system for new event detection. In *Proceedings of the 26th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 330–337. ACM.
- Deepayan Chakrabarti and Kunal Punera. 2011. Event summarization using tweets. *ICWSM*, 11:66–73.
- Kaustav Das, Jeff Schneider, and Daniel B Neill. 2008. Anomaly pattern detection in categorical datasets. In *Proceedings of the 14th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 169–176. ACM.
- Anish Das Sarma, Alpa Jain, and Cong Yu. 2011. Dynamic relationship and event discovery. In *Proceedings of the fourth ACM international conference on Web search and data mining*, pages 207–216. ACM.
- Wenwen Dou, Xiaoyu Wang, Drew Skau, William Ribarsky, and Michelle X Zhou. 2012. Leadline: Interactive visual analysis of text data through event identification and exploration. In *Visual Analytics Science and Technology (VAST), 2012 IEEE Conference on*, pages 93–102. IEEE.
- Gabriel Pui Cheong Fung, Jeffrey Xu Yu, Philip S Yu, and Hongjun Lu. 2005. Parameter free bursty events detection in text streams. In *Proceedings of the 31st international conference on Very large data bases*, pages 181–192. VLDB Endowment.
- Wei Gao, Peng Li, and Kareem Darwish. 2012. Joint topic modeling for event summarization across news and social media streams. In *Proceedings of the 21st ACM international conference on Information and knowledge management*, pages 1173–1182. ACM.
- Zoubin Ghahramani and Matthew J Beal. 2001. Propagation algorithms for variational bayesian learning. *Advances in neural information processing systems*, pages 507–513.
- Prem K Gopalan, Laurent Charlin, and David Blei. 2014. Content-based recommendations with poisson factorization. In Z. Ghahramani, M. Welling, C. Cortes, N.D. Lawrence, and K.Q. Weinberger, editors, *Advances in Neural Information Processing Systems 27*, pages 3176–3184. Curran Associates, Inc.
- Alan Jackoway, Hanan Samet, and Jagan Sankaranarayanan. 2011. Identification of live news events using twitter. In *Proceedings of the 3rd ACM SIGSPATIAL International Workshop on Location-Based Social Networks*, pages 25–32. ACM.
- Giridhar Kumaran and James Allan. 2004. Text classification and named entities for new event detection. In *Proceedings of the 27th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 297–304. ACM.
- Jey Han Lau, Nigel Collier, and Timothy Baldwin. 2012. On-line trend analysis with topic models: \# twitter trends detection topic model online. In *COLING*, pages 1519–1534.
- Zhiwei Li, Bin Wang, Mingjing Li, and Wei-Ying Ma. 2005. A probabilistic model for retrospective news event detection. In *Proceedings of the 28th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 106–113. ACM.
- Scott W Linderman and Ryan P Adams. 2014. Discovering latent network structure in point process data. *arXiv preprint arXiv:1402.0914*.
- Xueliang Liu, Raphaël Troncy, and Benoit Huet. 2011. Using social media to identify events. In *Proceedings of the 3rd ACM SIGMM international workshop on Social media*, pages 3–8. ACM.
- Michael Mathioudakis, Nilesh Bansal, and Nick Koudas. 2010. Identifying, attributing and describing spatial bursts. *Proceedings of the VLDB Endowment*, 3(1-2):1091–1102.
- Daniel B Neill, Andrew W Moore, Maheshkumar Sabhnani, and Kenny Daniel. 2005. Detection of emerging space-time clusters. In *Proceedings of the eleventh ACM SIGKDD international conference on Knowledge discovery in data mining*, pages 218–227. ACM.
- Wei Peng, Charles Perng, Tao Li, and Haixun Wang. 2007. Event summarization for system management. In *Proceedings of the 13th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 1028–1032. ACM.
- Timo Reuter and Philipp Cimiano. 2012. Event-based classification of social media streams. In *Proceedings of the 2nd ACM International Conference on Multimedia Retrieval*, page 22. ACM.
- Takeshi Sakaki, Makoto Okazaki, and Yutaka Matsuo. 2010. Earthquake shakes twitter users: real-time event detection by social sensors. In *Proceedings of the 19th international conference on World wide web*, pages 851–860. ACM.
- Hassan Sayyadi, Matthew Hurst, and Alexey Maykov. 2009. Event detection and tracking in social streams. In *ICWSM*.
- Aaron Schein, John Paisley, David M Blei, and Hanna Wallach. 2015. Bayesian poisson tensor factorization for inferring multilateral relations from sparse dyadic event counts. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1045–1054. ACM.
- Courtland VanDam. 2012. A probabilistic topic modeling approach for event detection in social media. Master’s thesis, Michigan State University.

768	Martin J. Wainwright and Michael I. Jordan. 2008. Graphical models, exponential families, and variational inference. <i>Found. Trends Mach. Learn.</i> , 1(1-2):1–305, January.	816
769		817
770		818
771		819
772	Xuanhui Wang, ChengXiang Zhai, Xiao Hu, and Richard Sproat. 2007. Mining correlated bursty topic patterns from coordinated text streams. In <i>Proceedings of the 13th ACM SIGKDD international conference on Knowledge discovery and data mining</i> , pages 784–793. ACM.	820
773		821
774		822
775		823
776	Gary M Weiss and Haym Hirsh. 1998. Learning to predict rare events in event sequences. In <i>KDD</i> , pages 359–363.	824
777		825
778		826
779	Yi Zhang, Jamie Callan, and Thomas Minka. 2002. Novelty and redundancy detection in adaptive filtering. In <i>Proceedings of the 25th annual international ACM SIGIR conference on Research and development in information retrieval</i> , pages 81–88. ACM.	827
780		828
781		829
782		830
783	Qiankun Zhao, Prasenjit Mitra, and Bi Chen. 2007. Temporal and information flow based event detection from social text streams. In <i>AAAI</i> , volume 7, pages 1501–1506.	831
784		832
785		833
786		834
787	Wayne Xin Zhao, Rishan Chen, Kai Fan, Hongfei Yan, and Xiaoming Li. 2012. A novel burst-based text representation model for scalable event detection. In <i>Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics: Short Papers-Volume 2</i> , pages 43–47. Association for Computational Linguistics.	835
788		836
789		837
790		838
791		839
792		840
793		841
794		842
795		843
796		844
797		845
798		846
799		847
800		848
801		849
802		850
803		851
804		852
805		853
806		854
807		855
808		856
809		857
810		858
811		859
812		860
813		861
814		862
815		863