

Detecting and Characterizing Events

Anonymous EMNLP submission

Abstract

Significant events are characterized by interactions between entities (e.g., countries, organizations, individuals) that deviate from typical interaction patterns. Investigators, such as historians, commonly read large quantities of text to construct an accurate picture of who, what, when, and where and event happened. In this work, we present the Capsule model for analyzing documents to identify and characterize events of potential significance. Specifically, we develop a model based on topic modeling to distinguish between topics that describe “business-as-usual” and topics that deviate from these patterns. To demonstrate this model, we analyze a corpus of over 2 million US State Department cables from the 1970s; we provide open-source implementations of an inference algorithm for the Capsule model and a pipeline to explore its results.

1 Introduction

Foreign embassies of the United States government communicate with each other and with the U.S. State Department through cabled message. The National Archive collects these documents in a running corpus, which traces the (unclassified) diplomatic history of the United States. Between 1973 and 1978, for example, it has collected about two million cables.

Typically, a cable from this collection describes diplomatic “business as usual,” such as arrangements for visiting officials, recovery of lost or stolen passports, or obtaining lists of names for meetings and conferences. For example, the embassies sent 8,635

cables during the week of April 21, 1975. Here is one, selected at random,

Hoffman, UNESCO Secretariat, requested info from PermDel concerning an official invitation from the USG RE subject meeting scheduled 10-13 JUNE 1975, Madison, Wisconsin. Would appreciate info RE status of action to be taken in order to inform Secretariat. Hoffman communicating with Dr. John P. Klus RE list of persons to be invited.

But hidden in the corpus are also cables about important diplomatic events, the cables and events that are of primary interest to historians. During that same week the United States was in the last moments of the Vietnam war and, on April 30, 1975, lost its hold on Saigon. This resulted in the end of the Vietnam War and a mass exodus of refugees from the country. One of the cables around this event is

GOA program to move Vietnamese Refugees to Australia is making little progress and probably will not cover more than 100-200 persons. Press comment on smallness of program has recognized difficulty of getting Vietnamese out of Saigon, but “Canberra Times” Apr 25 sharply critical of government’s performance. [...] Labor government clearly hopes whole matter will somehow disappear.

Our goal in this paper is to develop a method to help historians and political scientists wade through their collections, such as the 1970s cables, to find potentially important events, such as the fall of Saigon,




Figure 1: Measure of time interval impact on cable content: $\psi_t \frac{1}{|D_t|} \sum_{d \in D_t} \epsilon_{d,t}$, where D_t is the set of all cables sent in interval t . The grey background indicates the number of cables sent over time.

and the primary sources around them. We develop *Capsule*, a probabilistic model for detecting and characterizing important events in large collections of historical communication.

Figure 1 illustrates Capsule’s analysis of the two million cables from the National Archives. The y-axis is “eventness”, a loose measure how strongly a week’s cables deviate from the usual diplomatic chatter to discuss a matter that is common to many embassies. (This is described in detail in ??.)

The figure shows that Capsule detects many of the important moments during this five-year span, including Indonesia’s invasion of East Timor (XXX), the Air France hijacking and Israeli rescue operation (XXX), and the fall of Saigon (XXX). It also identifies other moments, such as the U.S. sharing lunar rocks with other countries (XXX) and the death of Mao Tse-tung (XXX). Broadly speaking, Capsule gives a picture of the diplomatic history of these five years; it identifies and characterizes moments and source material that might be of interest to a historian.

The intuition behind Capsule is this. Embassies write cables throughout the year, usually describing typical business such as the visiting of a government official. Sometimes, however, there is an important event—e.g., the fall of Saigon. When an event occurs, it pulls embassies away from their typical business to write cables that discuss what happened and its

consequences. Thus Capsule effectively defines an “event” to be a moment in history when embassies deviate from what each usually discusses, and when each embassy deviates in the same way.

Capsule embeds this intuition into a Bayesian model. It uses hidden variables to encode what “typical business” means for each embassy, how to characterize the events of each week, and which cables discuss those events. Given a corpus, the corresponding posterior distribution provides a filter on the cables that isolates important moments in the diplomatic history. Figure 1 illustrates this posterior.

Related work. We first review previous work on automatic event detection and other related concepts.

In both univariate and multivariate settings, the goal is often the same: analysts want to predict whether or not a rare events will occur (Weiss and Hirsh, 1998; Das et al., 2008). Capsule, in contrast, is designed to help analysts explore and understand the original data: our goal is interpretability, not prediction.

A common goal is to identify clusters of documents; these approaches are used on news articles (Zhao et al., 2012; Zhao et al., 2007; Zhang et al., 2002; Li et al., 2005; Wang et al., 2007; Allan et al., 1998) and social media posts (VanDam, 2012; Lau et al., 2012; Jackoway et al., 2011; Sakaki et al., 2010; Reuter and Cimiano, 2012; Becker et al., 2010; Sayyadi et al., 2009). In the case of news

articles, the task is to create new clusters as novel news stories appear—this does not help disentangle typical content from rare events of interest. Social media approaches identify rare events, but the methods are designed for short, noisy documents; they are not appropriate for larger documents that contain information about a variety of subjects.

Many existing methods use document terms as features, frequently weighted by tf-idf value (Fung et al., 2005; Kumaran and Allan, 2004; Brants et al., 2003; Das Sarma et al., 2011; Zhao et al., 2007; Zhao et al., 2012); here, events are bursts in groups of terms.

Topic models (Blei, 2012) reduce the dimensionality of text data; they have been used to help detect events mentioned in social media posts (Lau et al., 2012; Dou et al., 2012) and posts relevant to monitored events (VanDam, 2012). We rely on topic models to characterize both typical content and events, but grouped observations can also be summarized directly (Peng et al., 2007; Chakrabarti and Punera, 2011; Gao et al., 2012).

In addition to text data over time, author (Zhao et al., 2007), news outlet (Wang et al., 2007), and spatial information (Neill et al., 2005; Mathioudakis et al., 2010; Liu et al., 2011) can be used to augment event detection. Capsule uses author information in order to characterize typical concerns of authors.

Detecting and characterizing relationships (Schein et al., 2015; Linderman and Adams, 2014; Das Sarma et al., 2011) is related to event detection. When a message recipient is known, Capsule’s author input can be replaced with a sender-receiver pair, but the model could be further tailored for interactions within networks.

2 The Capsule Model

In this section we develop the Capsule model. Capsule captures patterns in entity behavior and identifies time intervals in which many entities deviate from these patterns. The model relies on high dimensional entity behavior data over time, such as text messages being sent between entities. We first review topic models at a high level, and give the intuition on Capsule. Then, we formally specify the model, discuss its relationship to Poisson processes, describe how we learn the hidden variables.

Background: Topic Models. Capsule builds on topic models. Topic models are algorithms for discovering the main themes in a large collection of documents; each document can then be summarized in terms of the global themes. More formally, a topic k is a probability distribution over the set of vocabulary words. Each document d is represented as a distribution over topics θ_d . Thus we can imagine that when we generate a document, we first pick which topics are relevant (and in what proportions). Under the LDA topic model (Blei et al., 2003), we know the number of words in each document. Then, for each word, we select a single topic from this distribution over topics, and finally select a vocabulary term from the corresponding topic’s distribution over the vocabulary. Alternatively, we can cast topic modeling as factorization, such as in Poisson factorization (Gopalan et al., 2014), and draw a word count for each term in the vocabulary.

The Capsule Model. Topic models are often applied to provide a structure for an otherwise unstructured collection of documents. Documents, however, are often accompanied by metadata, such as the date written or author attribution; this information is not exploited by traditional topic models. The Capsule model uses both author and date information to identify and characterize events that influence the content of the collection.

Consider an entity like the Bangkok American embassy, shown in Figure 2. We can imagine that there is a stream of messages (or *diplomatic cables*) being sent by this embassy—some might be sent to the US State Department, others to another American embassy like Hong Kong. An entity will usually talk about certain topics; the Bangkok embassy, for instance, is concerned with topics regarding southeast Asia more generally.

Now imagine that at a particular time t , an event occurs, such as the capture of Saigon during the Vietnam war. We do not directly observe that events occur, but we do observe the message stream. Using this stream, each event be described as a distribution over the vocabulary, similar to how topics are distributions over these same terms. When an event occurs, the message content changes for multiple entities. The day following the capture of Saigon, the majority of the diplomatic cables sent by the Bangkok embassy were about Vietnam war refugees. Thus we

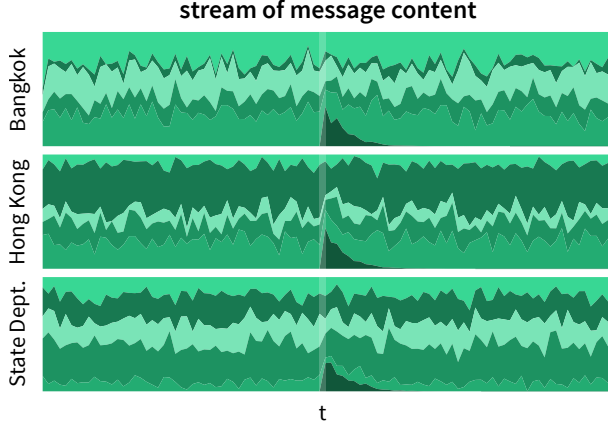


Figure 2: Cartoon intuition of Capsule; the y axis is the stacked proportion of messages about various subjects during a given time interval. The Bangkok embassy, Hong Kong embassy, and State Department all have typical concerns about which they usually send messages. When an event occurs at time t , the stream of message content alters to include the event, then fades back to “business as usual.” Capsule discovers both entities’ typical concerns and the event locations and content.

imagine that an entity’s stream of messages is controlled by what it usually talks about as well as the higher level stream of unobserved events.

Model Specification. We formally describe Capsule. The observed data are word counts $w_{d,v}$ for document d and vocabulary term v ; each document d also has an author (or entity) a_d and a time (or date) interval i_d associated with it.

The hidden variables of this model are general topics of conversation β , authors’ typical concerns ϕ , event descriptions π , event strengths ψ , and document-specific topics θ and event relevancy ϵ .

As in topic modeling, we represent the general topics of conversation with a $K \times V$ matrix β , where K is a low dimensional number of topics that we wish to capture, and V is the size of our vocabulary; each row β_k is normalized such that it represents the probability of seeing vocabulary word v when discussing topic k . As a generative process, we draw these general topics from a Dirichlet distribution, or $\beta_k \sim \text{Dirichlet}_V(\alpha_\beta)$.

In addition to using these general topics to represent entity concerns, each entity n has its own exclusive topic $\beta_0^{(n)}$, which can be appended as a bias row to the general topics β . These entity-specific topics are drawn from a Dirichlet, just as the general topics.

The concerns of author n are represented with ϕ_n , a $(K + 1)$ -dimensional topic vector, where each element is drawn from a gamma distribution, or $\phi_{n,k} \sim \text{Gamma}(s_\phi, r_\phi)$,¹ and the first element of the concern vector $\phi_{n,0}$ describes how much the entity n relies on its exclusive topic $\beta_0^{(n)}$.

Similar to topic modeling, we represent the contents of each document in topic space; each document d has a $(K + 1)$ -dimensional latent parameter θ_d to describe the particular contents of that document. Unlike traditional topic models, each document d ’s topics depend on the concerns of the author a_d ; each document topic $\theta_{d,k}$ is drawn from a gamma distribution parameterized by the corresponding author concerns $\phi_{a_d,k}$: $\theta_{d,k} \sim \text{Gamma}(s_\theta, \phi_{a_d,k})$.

To represent events, we consider discrete intervals of time. Each interval t has a corresponding interval strength ψ_t and description π_t . Event strengths are a single value for each interval t , and are drawn from a gamma distribution: $\psi_{n,k} \sim \text{Gamma}(s_\psi, r_\psi)$. These strengths indicate how important the interval is in determining message content. Interval descriptions are similar to topics: each description is a V -dimensional vector drawn from a Dirichlet distribution over the vocabulary terms, or $\pi_k \sim \text{Dirichlet}_V(\alpha_\pi)$.

Just as we describe each document d in terms of relevant topics with the θ_d parameters, we also describe the relevancy of each time interval with the ϵ_d parameters. These interval relevancy parameters are drawn from gamma distributions and depend on the overall strength ψ of the corresponding interval; for interval t and document d (written at time i_d), we have $\epsilon_{d,t} \sim \text{Gamma}(s_\epsilon, \psi_{i_d,t})$.

Conditional on the hidden variables and the author and time metadata, Capsule is a model of how document word counts came to be. For document d and vocabulary term v , we generate the word counts form a Poisson distribution parameterized by the documents topics θ_d and relevant events ϵ , as well as global topic β and event descriptions π :

$$w_{d,v} \sim \text{Poisson} \left(\theta_d^\top \beta_v^{(a_d)} + \sum_{t=1}^T f(i_d, t) \epsilon_{d,t} \pi_{t,v} \right), \quad (1)$$

¹We use the shape-rate parameterization for all Gamma distributions.

where f is some function of decay. This function is important because events should not remain at their full strength indefinitely, but should decay over time. In our experiments, we consider step functions, linear decay, and exponential decay. Figure 3 gives the full generative process for Capsule.

- for each time step $t = 1:T$,
 - draw interval description over vocabulary $\pi_t \sim \text{Dirichlet}_V(\alpha)$
 - draw interval strength $\psi_t \sim \text{Gamma}(s_\psi, r_\psi)$
- for each entity $n = 1:N$,
 - draw entity-specific topics over vocabulary $\beta_0^{(n)} \sim \text{Dirichlet}_V(\alpha)$
 - draw entity-specific topic strength $\phi_{n,0} \sim \text{Gamma}(s_\phi, r_\phi)$
- for each topic $k = 1:K$,
 - draw general topic distribution over vocabulary $\beta_k \sim \text{Dirichlet}_V(\alpha)$
 - for each entity $n = 1:N$,
 - ▶ draw general entity concern $\phi_{n,k} \sim \text{Gamma}(s_\phi, r_\phi)$
- for each document $d = 1:D$ sent at time i_d by author a_d ,
 - draw local entity concern $\theta_{d,0} \sim \text{Gamma}(s_\theta, \phi_{a_d,0})$
 - for each topic $k = 1:K$,
 - ▶ draw local entity concern $\theta_{d,k} \sim \text{Gamma}(s_\theta, \phi_{a_d,k})$
 - for each time $t = 1:T$,
 - ▶ draw local interval relevancy $\epsilon_{d,t} \sim \text{Gamma}(s_\epsilon, \psi_{i_d,t})$
 - for each vocabulary term $v = 1:V$,
 - ▶ draw word counts $w_{d,v} \sim \text{Poisson}\left(\theta_d^\top \beta_v^{(a_d)} + \sum_{t=1}^T f(i_d, t) \epsilon_{d,t} \pi_{t,v}\right)$

Figure 3: The generative process for Capsule.

Relationship to Poisson Processes. The Capsule model includes a specific variety of Poisson process. Poisson processes describe the number of discrete observations between times a and b as being drawn from a Poisson distribution parameterized by the in-

tegral of some intensity function $\lambda(t)$, or

$$N(a, b] \sim \text{Poisson}\left(\int_a^b \lambda(t) dt\right).$$

In the case of Capsule, we have a Poisson process for every combination of document d and vocabulary term v , which generate our observed word counts w .

This collection of Poisson processes have a base rate for each intensity function; this captures the “business-as-usual” content which is described by general and entity topics β and document-specific concerns θ . The intensity functions λ also have an excitatory component, which are influenced by *external events*—in the case of National Archive cables, we interpret these as real-world historical events. This excitatory aspect is modeled by the time interval relevancy parameters ϵ , interval descriptions π , and decay function f .

Similar to existing work on network influence that uses Hawkes processes (Linderman and Adams, 2015; Guo et al., 2014), Capsule assumes discrete time intervals for both the observations and the external events. Note that while the model is excitatory, it is not self or mutually exciting like the network models. Instead, the events that cause excitation are not the observations w , but external events modeled by Capsule. Capsule assumes that only one event can occur in each time interval t , and that it is characterized by its description π_t and strength ψ_t .

Learning the hidden variables. In order to use the Capsule model to explore the observed documents, we must compute the posterior distribution. Conditional on the observed word counts w , our goals to to compute the posterior values of the hidden parameters—global interval strengths ψ , interval descriptions π , entity concerns ϕ , and topics β , as well as document-specific entity concerns θ and interval relevancy parameters ϵ .

As for many Bayesian models, the exact posterior for Capsule is not tractable to compute; approximating it is our central statistical and computational problem. We develop an approximate inference algorithm for Capsule based on variational methods (Wainwright and Jordan, 2008),² which is detailed in Appendix A. This algorithm produces a fitted variational distribution which can then be used as a proxy for

²Source code is available at <https://github.com/????/capsule>.

the true posterior, allowing us to explore a collection of documents with Capsule.

Detecting and characterizing events. pgh here... footnote on visualization

3 Evaluation

In this section we explore the performance of Capsule on a collection of US State Department cables. These cables were sent between 1973 and 1978 and obtained from the History Lab at Columbia,³ which received them from the Central Foreign Policy Files at the National Archives. In addition to the text of the cables themselves, each document is supplemented with information about who sent the cable (e.g., the State Department, the U.S. Embassy in Saigon, or an individual by name), who received the cable (often multiple entities), and the date the cable was sent. To test our model, we used a vocabulary of size 6,293 and omitted cables with fewer than three terms, resulting in a collection of 2,139,324 messages sent between 27,134 entities. We selected a weekly duration for the time intervals, as few cables were sent on the weekends.

We fit Capsule with 100 topics and using an exponential decay with mean lifetime of 3—this indicates that most intervals would no longer be relevant after about 3 weeks. To detect when an event occurs, we multiply the average event relevancy ϵ for all documents in a given interval together with interval strength ψ , or $\psi_t \frac{1}{|D_t|} \sum_{d \in D_t} \epsilon_{d,t}$, where D_t is the set of all cables sent in interval t .

Figure 1 shows this measure over the duration of the data set. The highest time intervals, ones in which we declare events to be detected, include the tallest peak the week of December 1, 1975, just prior to the Indonesian invasion of East Timor, which Began December 7, 1975. The second tallest peak occurs the week of April 21, 1975, just prior to the fall of Saigon on April 30, 1975. For any given week, we can sort the documents by their interval relevancy parameters ϵ ; Tables 1 and 2 show the top cables for these two events, which reflect the real-world events those weeks.

Other event peaks include the week of July 2, 1973; the top three words under event its description π are *bicentennial*, *hijack*, and *mercenary*. Top cables

³<http://history-lab.org>

under event relevancy ϵ surround the bicentennial celebration of United States (July 4, 1973) and the Air France hijacking incident that began on June 27: Israeli operatives rescued hostages from this incident on July 4th.

Another peak occurs the week of April 17, 1978 surrounding a UN special session on disarmament; the top three words under event its description π are *SSOD* (acronym for “special session on disarmament”, *disarmament*, and *ICS* (likely an acronym for “incident command system”).

Examples of general topics of conversation are shown in Table 3 and entity-exclusive topics are shown in Table 4; these show us how entity topics absorb location-specific words, preventing these terms from overwhelming the general topics.

These exploratory results show that our model is successfully capturing when multiple entities are discussing the same subjects and that our model can be used to explore the underlying data by providing a structured scaffold from which to view the data.

We also considered held out log likelihood in evaluating the fitness of our model, as shown in Table 5. We see that the event structure is crucial to fitting the data well. We held out 5% of the data.




Figure 4: Event detection performance on twenty simulated datasets. Capsule is able to detect events as well as comparison methods, but its performance has higher variance.

Finally, we simulated data according to our generative process in order to compare our method to baseline and existing approaches. To evaluate event detection, we created a ranked list of all time intervals and computed the overlap between a model and the simulated ground at every threshold; this generates an curve under which we can compute the area and normalized based on ideal performance—we refer to this metric as event detection AUC. The

ϵ	date	entity	subject
0.1237	1975-12-03	STATE	PRESIDENT'S TALKING POINTS ON PORTUGUESE TIMOR
0.1210	1975-12-03	STATE	PRESIDENT'S TALKING POINTS ON PORTUGUESE TIMOR
0.1153	1975-12-04	STATE	TIMOR WE ARE REPEATING FYI A DAO MESSAGE
0.1126	1975-012-04	STATE	LEGAL PROBLEMS RELATING TO PORTUGUESE TIMOR
0.1053	1975-12-07	SECRETARY PEKING	US SUPPORT FOR TIMOR RESOLUTION
0.1021	1975-12-01	STATE	INVASION OF PORTUGUESE TIMOR

Table 1: Top documents for the time interval of week December 1, 1975, just prior to the Indonesian invasion of East Timor, which Began December 7, 1975.

ϵ	date	entity	subject
0.0908	1975-04-24	MANSFIELD, MIKE	ASSISTANCE IN EVACUATING FAMILY FROM SOUTH VIETNAM
0.0886	1975-04-24	RAILSBACK, TOM	ASSISTANCE IN EVACUATING FRIEND FROM SOUTH VIETNAM
0.0877	1975-04-24	MANSFIELD, MIKE	ASSISTANCE IN EVACUATING FAMILY FROM SOUTH VIETNAM
0.0863	1975-04-24	WILLIAMS, HARRISON	ASSISTANCE IN EVACUATING FAMILY FROM VIETNAM
0.0860	1975-04-24	KOCH, EDWARD	ASSISTANCE IN EVACUATING FAMILY FROM SOUTH VIETNAM
0.0858	1975-04-21	SCHWEIKER, RICHARD	SUPPORT IN EVACUATING FAMILY FROM VIETNAM
\vdots	\vdots	\vdots	\vdots
0.0812	1975-04-25	KETCHUM, WILLIAM	MOVEMENT OF SOUTH VIETNAMESE REFUGEES TO GUAM
0.0800	1975-04-21	SCOTT, HUGH	WHEREABOUTS OF MISSIONARIES IN VIETNAM

Table 2: Top documents for the time interval of week April 21, 1975, just prior to the fall of Saigon on April 30, 1975.

most successful of the comparison methods for event detection was average absolute error in wordcount, both unweighted and weighted by tf-idf. Figure 4 shows that Capsule can outperform these approaches for event detection, but that it has higher variance in performance. The other comparison method in Figure 4 is based on LDA; we fit a multinomial Gaussian to the topic representation of all documents and then computed the average probability of seeing the topic distributions of documents in the time interval. Time intervals with the lowest probability were marked as most likely to have events. All other baselines performed close to random for event detection.

This method of fitting a multinomial Gaussian to LDA representations of documents also performed well for recovering relevant documents. This approach can be altered to fit a per-entity multinomial Gaussian, but this performs worse. Simply finding documents based on absolute deviation from the mean works well in LDA topic space (relative to overall mean or entity mean), but not over the full vocabulary. Word count deviations, which performed well for event detection, performed worse than random for document recovery. Both Capsule and its




Figure 5: Precision of recovering the top ten most relevant documents, averaged over all time intervals. Capsule performs best, averaged over twenty simulations.

event-only partial model outperform all comparison methods in terms of document recovery. Figure 5 shows precision of recovering the top ten documents.

We assessed the sensitivity of our model to three different decay functions f : exponential, linear, and step functions. We simulated data for each function and then fit Capsule using every permutation of f and multiple settings for event decay duration. In all cases, we found that the model is not sensitive to decay shape or duration.

top terms
OUTLOOK, REVIEW, HIRE, PERSONNEL, INVITE, PREPARE, NECESSARY ARREST, INCIDENT, SECURITY, FAMILY, OFF, GUARD, DEATH, JAIL LOCATE, HOME, SON, DEATH, PLEASE, CONTACT, FATHER, DEPARTMENT REQUEST, REFUGEE, RESPONSE, SERVICE, SALE, ASYLUM, APPRECIATE MARKET, REPORT, COPY, COMMERCIAL, FOOD, IMPORT, COMMERCE FEAR, LEADERSHIP, BACK, ARM, ROLE, PLAY, THREATEN HOTEL, TRAVEL, RESERVATION, VISIT, ARRANGE, SCHEDULE, STAY

Table 3: Top vocabulary terms for a selection of topics, according to topic distributions β_k .

entity	top terms
STATE	REQUEST, FOLLOW, EMBASSY, MEET, MAKE, STATE, DEPARTMENT
BANGKOK	BANGKOK, THAILAND, THAI, REFUGEE, EMBASSY, FOLLOW, REPORT
JERUSALEM	JERUSALEM, ISRAELI, BANK, REPORT, SAY, COMMENT, ONE
STOCKHOLM	SWEDISH, SWEDEN, TRADE, MEET, EMBASSY, FOLLOW, MAKE
CASABLANCA	CASABLANCA, MOROCCO, MOROCCAN, REQUEST, PLEASE, FOLLOW, NOTE
KAMPALA	UGANDAN, NAIROBI, AFRICAN, IMPERIALIST, VOICE, KENYA, MISSIONARY
NDJAMEAN	CHADIAN, CHAD, LAGOS, DROUGHT, INITIATION, AUSTERITY, GOC

Table 4: Top vocabulary terms for a selection of entities according to entity-exclusive topics $\beta_0^{(n)}$.

method	held out log likelihood
full Capsule model	-2.41e7
event only	-1.82e7
topics only	-2.69e7
entities only	-2.43e7

Table 5: Held-out data log likelihood.

4 Discussion

We have presented Capsule, a Bayesian model that identifies when events occur, characterizes these events, and discovers the typical concerns of author entities. We have shown that Capsule outperforms comparison methods and explored its results on a real-world datasets. We anticipate that Capsule can be used by historians, political scientists, and others who wish to investigate events in large text corpora.

References

- James Allan, Ron Papka, and Victor Lavrenko. 1998. On-line new event detection and tracking. In *Proceedings of the 21st annual international ACM SIGIR conference on Research and development in information retrieval*, pages 37–45. ACM.
- Hila Becker, Mor Naaman, and Luis Gravano. 2010. Learning similarity metrics for event identification in

- social media. In *Proceedings of the third ACM international conference on Web search and data mining*, pages 291–300. ACM.
- David M. Blei, Andrew Y. Ng, and Michael I. Jordan. 2003. Latent Dirichlet allocation. *JMLR*, 3:993–1022, March.
- David M Blei. 2012. Probabilistic topic models. *Communications of the ACM*, 55(4):77–84.
- Thorsten Brants, Francine Chen, and Ayman Farahat. 2003. A system for new event detection. In *Proceedings of the 26th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 330–337. ACM.
- Deepayan Chakrabarti and Kunal Punera. 2011. Event summarization using tweets. *ICWSM*, 11:66–73.
- Kaustav Das, Jeff Schneider, and Daniel B Neill. 2008. Anomaly pattern detection in categorical datasets. In *Proceedings of the 14th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 169–176. ACM.
- Anish Das Sarma, Alpa Jain, and Cong Yu. 2011. Dynamic relationship and event discovery. In *Proceedings of the fourth ACM international conference on Web search and data mining*, pages 207–216. ACM.
- Wenwen Dou, Xiaoyu Wang, Drew Skau, William Ribarsky, and Michelle X Zhou. 2012. Leadline: Interactive visual analysis of text data through event identification and exploration. In *Visual Analytics Science and Technology (VAST), 2012 IEEE Conference on*, pages 93–102. IEEE.

- Gabriel Pui Cheong Fung, Jeffrey Xu Yu, Philip S Yu, and Hongjun Lu. 2005. Parameter free bursty events detection in text streams. In *Proceedings of the 31st international conference on Very large data bases*, pages 181–192. VLDB Endowment.
- Wei Gao, Peng Li, and Kareem Darwish. 2012. Joint topic modeling for event summarization across news and social media streams. In *Proceedings of the 21st ACM international conference on Information and knowledge management*, pages 1173–1182. ACM.
- Prem K Gopalan, Laurent Charlin, and David Blei. 2014. Content-based recommendations with poisson factorization. In Z. Ghahramani, M. Welling, C. Cortes, N.D. Lawrence, and K.Q. Weinberger, editors, *Advances in Neural Information Processing Systems 27*, pages 3176–3184. Curran Associates, Inc.
- Fangjian Guo, Charles Blundell, Hanna Wallach, and Katherine Heller. 2014. The bayesian echo chamber: Modeling social influence via linguistic accommodation. *arXiv preprint arXiv:1411.2674*.
- Alan Jackoway, Hanan Samet, and Jagan Sankaranarayanan. 2011. Identification of live news events using twitter. In *Proceedings of the 3rd ACM SIGSPATIAL International Workshop on Location-Based Social Networks*, pages 25–32. ACM.
- Giridhar Kumaran and James Allan. 2004. Text classification and named entities for new event detection. In *Proceedings of the 27th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 297–304. ACM.
- Jey Han Lau, Nigel Collier, and Timothy Baldwin. 2012. On-line trend analysis with topic models: \# twitter trends detection topic model online. In *COLING*, pages 1519–1534.
- Zhiwei Li, Bin Wang, Mingjing Li, and Wei-Ying Ma. 2005. A probabilistic model for retrospective news event detection. In *Proceedings of the 28th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 106–113. ACM.
- Scott W Linderman and Ryan P Adams. 2014. Discovering latent network structure in point process data. *arXiv preprint arXiv:1402.0914*.
- Scott W Linderman and Ryan P Adams. 2015. Scalable bayesian inference for excitatory point process networks. *arXiv preprint arXiv:1507.03228*.
- Xueliang Liu, Raphaël Troncy, and Benoit Huet. 2011. Using social media to identify events. In *Proceedings of the 3rd ACM SIGMM international workshop on Social media*, pages 3–8. ACM.
- Michael Mathioudakis, Nilesh Bansal, and Nick Koudas. 2010. Identifying, attributing and describing spatial bursts. *Proceedings of the VLDB Endowment*, 3(1-2):1091–1102.
- Daniel B Neill, Andrew W Moore, Maheshkumar Sabhnani, and Kenny Daniel. 2005. Detection of emerging space-time clusters. In *Proceedings of the eleventh ACM SIGKDD international conference on Knowledge discovery in data mining*, pages 218–227. ACM.
- Wei Peng, Charles Perng, Tao Li, and Haixun Wang. 2007. Event summarization for system management. In *Proceedings of the 13th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 1028–1032. ACM.
- Timo Reuter and Philipp Cimiano. 2012. Event-based classification of social media streams. In *Proceedings of the 2nd ACM International Conference on Multimedia Retrieval*, page 22. ACM.
- Takeshi Sakaki, Makoto Okazaki, and Yutaka Matsuo. 2010. Earthquake shakes twitter users: real-time event detection by social sensors. In *Proceedings of the 19th international conference on World wide web*, pages 851–860. ACM.
- Hassan Sayyadi, Matthew Hurst, and Alexey Maykov. 2009. Event detection and tracking in social streams. In *ICWSM*.
- Aaron Schein, John Paisley, David M Blei, and Hanna Wallach. 2015. Bayesian poisson tensor factorization for inferring multilateral relations from sparse dyadic event counts. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1045–1054. ACM.
- Courtland VanDam. 2012. A probabilistic topic modeling approach for event detection in social media. Master’s thesis, Michigan State University.
- Martin J. Wainwright and Michael I. Jordan. 2008. Graphical models, exponential families, and variational inference. *Found. Trends Mach. Learn.*, 1(1-2):1–305, January.
- Xuanhui Wang, ChengXiang Zhai, Xiao Hu, and Richard Sproat. 2007. Mining correlated bursty topic patterns from coordinated text streams. In *Proceedings of the 13th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 784–793. ACM.
- Gary M Weiss and Haym Hirsh. 1998. Learning to predict rare events in event sequences. In *KDD*, pages 359–363.
- Yi Zhang, Jamie Callan, and Thomas Minka. 2002. Novelty and redundancy detection in adaptive filtering. In *Proceedings of the 25th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 81–88. ACM.
- Qiankun Zhao, Prasenjit Mitra, and Bi Chen. 2007. Temporal and information flow based event detection from social text streams. In *AAAI*, volume 7, pages 1501–1506.
- Wayne Xin Zhao, Rishan Chen, Kai Fan, Hongfei Yan, and Xiaoming Li. 2012. A novel burst-based text

864	representation model for scalable event detection. In	912
865	<i>Proceedings of the 50th Annual Meeting of the Asso-</i>	913
866	<i>ciation for Computational Linguistics: Short Papers-</i>	914
867	<i>Volume 2</i> , pages 43–47. Association for Computational	915
868	Linguistics.	916
869		917
870		918
871		919
872		920
873		921
874		922
875		923
876		924
877		925
878		926
879		927
880		928
881		929
882		930
883		931
884		932
885		933
886		934
887		935
888		936
889		937
890		938
891		939
892		940
893		941
894		942
895		943
896		944
897		945
898		946
899		947
900		948
901		949
902		950
903		951
904		952
905		953
906		954
907		955
908		956
909		957
910		958
911		959