

Detecting and Characterizing Events: Appendices

Allison J. B. Chaney
Princeton University
achaney@cs.princeton.edu

Hanna Wallach
Microsoft Research
wallach@microsoft.com

Matthew Connelly
Columbia University
mjc96@columbia.edu

David M. Blei
Columbia University
david.blei@columbia.edu

A Inference

In this appendix, we describe the details of the approximate inference algorithm for Capsule.

Conditioned on the observed term counts— n_{dv} for vocabulary term v in message d ; collectively \mathbf{N} —our goal is to learn the posterior distribution of the latent variables. Each message is associated with an author entity a_d and a time interval t_d within which that messages was sent. The latent variables are the general topics β_1, \dots, β_K , the entity topics η_1, \dots, η_A , and the event topics $\gamma_1, \dots, \gamma_T$, as well as the message-specific strengths $\theta_1, \dots, \theta_D$, ζ_1, \dots, ζ_D , and $\epsilon_1, \dots, \epsilon_D$, the entity-specific strengths ϕ_1, \dots, ϕ_A and ξ_1, \dots, ξ_A , and the event strengths ψ_1, \dots, ψ_T . See figures 3 and 4 for the graphical model and generative process.

As for many Bayesian models, the posterior distribution is not tractable to compute; we must instead approximate it. We therefore introduce an approximate inference algorithm for Capsule, based on variational methods (Jordan et al., 1999; Wainwright and Jordan, 2008). Variational methods approximate the true posterior distribution p with a (simpler) variational distribution q . Inference then consists of minimizing the KL divergence from q to p . This is equivalent to maximizing the evidence lower bound (ELBO):

$$\mathcal{L}(q) = \mathbb{E}_q [\log p(\mathbf{N}, \beta, \eta, \gamma, \theta, \zeta, \epsilon, \phi, \xi, \psi) - \log q(\beta, \eta, \gamma, \theta, \zeta, \epsilon, \phi, \xi, \psi)]. \quad (10)$$

We define q using the mean field assumption:

$$\begin{aligned} q(\beta, \eta, \gamma, \theta, \zeta, \epsilon, \phi, \xi, \psi) = & \prod_{d=1}^D \left(q(\zeta_d | \lambda_d) \prod_{k=1}^K q(\theta_{dk} | \lambda_{dk}^\theta) \prod_{t=1}^T q(\epsilon_{dt} | \lambda_{dt}^\epsilon) \right) \times \\ & \prod_{k=1}^K \left(q(\beta_k | \lambda_k^\beta) \prod_{a=1}^A q(\phi_{ak} | \lambda_{ak}^\phi) \right) \prod_{a=1}^A \left(q(\eta_a | \lambda_a^\eta) q(\xi_a | \lambda_a^\xi) \right) \prod_{t=1}^T \left(q(\gamma_t | \lambda_t^\gamma) q(\psi_t | \lambda_t^\psi) \right) \end{aligned} \quad (11)$$

The variational distributions for the topics $q(\beta_k)$, $q(\eta_a)$, and $q(\gamma_t)$ are all Dirichlet distributions with free variational parameters λ_k^β , λ_a^η , and λ_t^γ , respectively. The variational distributions for the strengths $q(\theta_{dk})$, $q(\zeta_d)$, $q(\epsilon_{dt})$, $q(\phi_{ak})$, $q(\xi_a)$, and $q(\psi_t)$ are all gamma distributions with free variational parameters λ_{dk}^θ , λ_d^ζ , λ_{dt}^ϵ , λ_{ak}^ϕ , λ_a^ξ , and λ_t^ψ , respectively. Each of these parameters has two components: shape s and rate r .

The expectations under q , which we need to maximize the ELBO, have closed analytic forms. We therefore update each free variational parameter in turn, following a standard coordinate-ascent approach.

To obtain update equations for the free variational parameters, we introduce auxiliary latent variables:

$$z_{dkv}^{\mathcal{K}} \sim \text{Poisson}(\theta_{dk}\beta_{kv}) \quad (12)$$

$$z_{dv}^{\mathcal{A}} \sim \text{Poisson}(\zeta_d\eta_{av}) \quad (13)$$

$$z_{d tv}^{\mathcal{T}} \sim \text{Poisson}(f(t_d, t)\epsilon_{dt}\gamma_{tv}), \quad (14)$$

where the superscripts \mathcal{K} , \mathcal{A} , and \mathcal{T} indicate the general, entity, and event topics, respectively. When marginalized out, these variables—collectively \mathbf{z} —leave the model intact. Because the Poisson distribution has an additive property, the value of n_{dv} is completely determined by the values of these variables:

$$n_{dv} = \sum_{k=1}^K z_{dkv}^{\mathcal{K}} + z_{dv}^{\mathcal{A}} + \sum_{t=1}^T z_{d tv}^{\mathcal{T}}. \quad (15)$$

Coordinate-ascent variational inference depends on the conditional distribution of each latent variable given the values of the other latent variables and the data. We use $D(a)$ to denote the set of messages sent by entity a and $D(t)$ to denote the set of messages potentially affected by event t (e.g., all messages sent after time interval t , in the case of an exponential decay function). The conditional distributions are:

$$(\beta_k \mid \mathbf{N}, \mathbf{z}, \eta, \gamma, \theta, \zeta, \epsilon, \phi, \xi, \psi) \sim \text{Dirichlet}_V \left(\alpha + \sum_{d=1}^D z_{dk1}^{\mathcal{K}}, \dots, \alpha + \sum_{d=1}^D z_{dkV}^{\mathcal{K}} \right) \quad (16)$$

$$(\eta_a \mid \mathbf{N}, \mathbf{z}, \beta, \gamma, \theta, \zeta, \epsilon, \phi, \xi, \psi) \sim \text{Dirichlet}_V \left(\alpha + \sum_{d \in D(a)} z_{d1}^{\mathcal{A}}, \dots, \alpha + \sum_{d \in D(a)} z_{dV}^{\mathcal{A}} \right) \quad (17)$$

$$(\gamma_t \mid \mathbf{N}, \mathbf{z}, \beta, \eta, \theta, \zeta, \epsilon, \phi, \xi, \psi) \sim \text{Dirichlet}_V \left(\alpha + \sum_{d \in D(t)} z_{d1t}^{\mathcal{T}}, \dots, \alpha + \sum_{d \in D(t)} z_{dVt}^{\mathcal{T}} \right) \quad (18)$$

$$(\theta_{dk} \mid \mathbf{N}, \mathbf{z}, \beta, \eta, \gamma, \zeta, \epsilon, \phi, \xi, \psi) \sim \text{Gamma} \left(s + \sum_{v=1}^V z_{dkv}^{\mathcal{K}}, \phi_{a_d k} + \sum_{v=1}^V \beta_{kv} \right) \quad (19)$$

$$(\zeta_d \mid \mathbf{N}, \mathbf{z}, \beta, \eta, \gamma, \theta, \epsilon, \phi, \xi, \psi) \sim \text{Gamma} \left(s + \sum_{v=1}^V z_{dv}^{\mathcal{A}}, \xi_{a_d} + \sum_{v=1}^V \eta_{av} \right) \quad (20)$$

$$(\epsilon_{dt} \mid \mathbf{N}, \mathbf{z}, \beta, \eta, \gamma, \theta, \zeta, \phi, \xi, \psi) \sim \text{Gamma} \left(s + \sum_{v=1}^V z_{d tv}^{\mathcal{T}}, \psi_t + f(t_d, t) \sum_{v=1}^V \gamma_{tv} \right) \quad (21)$$

$$(\phi_{ak} \mid \mathbf{N}, \mathbf{z}, \beta, \eta, \gamma, \theta, \zeta, \epsilon, \xi, \psi) \sim \text{Gamma} \left(s + |D(a)|s, r + \sum_{d \in D(a)} \theta_{dk} \right) \quad (22)$$

$$(\xi_a \mid \mathbf{N}, \mathbf{z}, \beta, \eta, \gamma, \theta, \zeta, \epsilon, \phi, \psi) \sim \text{Gamma} \left(s + |D(a)|s, r + \sum_{d \in D(a)} \zeta_d \right) \quad (23)$$

$$(\psi_t \mid \mathbf{N}, \mathbf{z}, \beta, \eta, \gamma, \theta, \zeta, \epsilon, \phi, \xi) \sim \text{Gamma} \left(s + |D(t)|s, r + \sum_{d \in D(t)} \epsilon_{dt} \right). \quad (24)$$

The conditional distribution of the auxiliary latent variables is:

$$(\langle \mathbf{z}_{dv}^{\mathcal{K}}, \mathbf{z}_{dv}^{\mathcal{A}}, \mathbf{z}_{dv}^{\mathcal{T}} \rangle \mid \mathbf{N}, \boldsymbol{\beta}, \boldsymbol{\eta}, \boldsymbol{\gamma}, \boldsymbol{\theta}, \boldsymbol{\zeta}, \boldsymbol{\epsilon}, \boldsymbol{\phi}, \boldsymbol{\xi}, \boldsymbol{\psi}) \sim \text{Mult}(n_{dv}, \boldsymbol{\omega}_{dv}), \quad (25)$$

where

$$\boldsymbol{\omega}_{dv} \propto \langle \theta_{d1}\beta_{1v}, \dots, \theta_{dK}\beta_{Kv}, \zeta_d\eta_{adv}, f(t_d, 1)\epsilon_{d1}\gamma_{1v}, \dots, f(t_d, T)\epsilon_{dT}\gamma_{Tv} \rangle. \quad (26)$$

Given the conditional distributions, coordinate-ascent variational inference involves setting each free variational parameter to the expected value of the corresponding model parameter under the variational distribution. We provide pseudocode in [algorithm 1](#); we use $\boldsymbol{\lambda}$ to denote the entire set of free variational parameters and $V(d)$ to denote the set of vocabulary terms present in document d . Our approximate inference algorithm produces a fitted variational posterior distribution which can then be used as a proxy for the true posterior distribution. The source code is available online at <https://github.com/ajbc/capsule>.

B Additional Results

In this appendix, we provide additional results for Capsule. [Table 5](#) lists the highest-ranked messages for an event described in [section 5](#). [Tables 6](#) and [7](#) contain examples of general topics and entity topics, respectively.

B.1 Decay Functions

We assessed Capsule’s sensitivity to several decay functions. We considered an exponential function,

$$f(t_d, t) = \begin{cases} 0 & t \leq t_d < t + \tau \\ \exp\left(\frac{-(t_d - t)}{\tau/5}\right) & \text{otherwise,}^1 \end{cases} \quad (27)$$

a linear function,

$$f(t_d, t) = \begin{cases} 1 - \frac{(t_d - t)}{\tau} & t \leq t_d < t + \tau \\ 0 & \text{otherwise,} \end{cases} \quad (28)$$

and a step function,

$$f(t_d, t) = \begin{cases} 1 & t \leq t_d < t + \tau \\ 0 & \text{otherwise.} \end{cases} \quad (29)$$

For each decay function, we used Capsule’s generative process, with $\tau = 3$, to create ten simulated data sets. We used three versions of our approximate inference algorithm—one for each decay function—to obtain a fitted variational distribution for each data set. We experimented with $\tau = 1, \dots, 5$. [Figure 7](#) shows event detection results (using [equation \(8\)](#)) and message identification results (using precision at ten messages).

As expected, Capsule performs best when the decay function used in the inference algorithm matches the decay function used to generate the data. For both event detection and message identification, the exponential function is least sensitive to the value of τ used to generate the data and the value of τ used in the inference algorithm. We also found that the exponential function gave the most interpretable results for real messages.

¹Unlike the linear and step functions, the exponential function could be evaluated for any time interval t after message d ’s appearance at t_d ; however, we truncate the function for computational reasons. The mean lifetime of the exponential decay is τ divided by five, which ensures that 99.3% of the area under the curve has occurred before we truncate the function at τ .

Algorithm 1: Coordinate-ascent variational inference for Capsule.

Input: observed term counts \mathbf{N}

Output: approximate posterior distribution of the latent variables, in terms of free variational parameters λ

Initialize $\mathbb{E}[\beta_k]$ to slightly random around uniform for each general topic k

Initialize $\mathbb{E}[\text{all other latent variables}]$ to uniform

for iteration $m = 1, \dots, M$ **do**

set $\lambda^{\theta,r}, \lambda^{\xi,r}$, and $\lambda^{\epsilon,r}$ to 0 and set remaining λ using priors

update $\lambda_{dk}^{\theta,r} += \sum_V \mathbb{E}[\beta_{kv}]$ for each message d and general topic k

for message $d = 1, \dots, D$ **do**

for term $v \in V(d)$ **do**

set ω_{dv} using expected values of the latent variables (equation (26))

set $\mathbb{E}[\langle \mathbf{z}_{dv}^{\mathcal{K}}, \mathbf{z}_{dv}^{\mathcal{A}}, \mathbf{z}_{dv}^{\mathcal{T}} \rangle] = n_{dv} \omega_{dv}$

update $\lambda_{kv}^{\beta} += \mathbb{E}[z_{dkv}^{\mathcal{K}}]$ for each general topic k (equation (16))

update $\lambda_{av}^{\eta} += \mathbb{E}[z_{dv}^{\mathcal{A}}]$ (equation (17))

update $\lambda_{tv}^{\gamma} += \mathbb{E}[z_{dtv}^{\mathcal{T}}]$ for each time interval t (equation (18))

update $\lambda_{dk}^{\theta,s} += \mathbb{E}[z_{dkv}^{\mathcal{K}}]$ for each general topic k (equation (19))

update $\lambda_d^{\xi,s} += \mathbb{E}[z_{dv}^{\mathcal{A}}]$ (equation (20))

update $\lambda_{dt}^{\epsilon,s} += \mathbb{E}[z_{dtv}^{\mathcal{K}}]$ for each time interval t (equation (21))

end

set $\lambda_{dk}^{\theta,r} = \mathbb{E}[\phi_{adk}] + \sum_v \mathbb{E}[\beta_{kv}]$ for each general topic k (equation (19))

set $\lambda_d^{\xi,r} = \mathbb{E}[\xi_{ad}] + \sum_v \mathbb{E}[\eta_{adv}]$ (equation (20))

set $\lambda_{dt}^{\epsilon,r} = \mathbb{E}[\psi_t] + f(t_d, t) \sum_v \mathbb{E}[\gamma_{tv}]$ for each time interval t (equation (21))

set $\mathbb{E}[\theta_{dk}] = \lambda_{dk}^{\theta,s} / \lambda_{dk}^{\theta,r}$ for each general topic k

set $\mathbb{E}[\xi_d] = \lambda_d^{\xi,s} / \lambda_d^{\xi,r}$

set $\mathbb{E}[\epsilon_{dt}] = \lambda_{dt}^{\epsilon,s} / \lambda_{dt}^{\epsilon,r}$ for each time interval t

update $\lambda_{adk}^{\phi,s} += s$ for each general topic k (equation (22))

update $\lambda_{ad}^{\xi,s} += s$ (equation (23))

update $\lambda_t^{\psi,s} += s$ for each time interval t where $f(t_d, t) \neq 0$ (equation (24))

update $\lambda_{adk}^{\phi,r} += \theta_{dk}$ for each general topic k (equation (22))

update $\lambda_{ad}^{\xi,r} += \xi_d$ (equation (23))

update $\lambda_t^{\psi,r} += \epsilon_{dt}$ for each time interval t (equation (24))

end

set $\mathbb{E}[\beta_k] = \lambda_k^{\beta} / \sum_v \lambda_{kv}^{\beta}$ for each general topic k

set $\mathbb{E}[\eta_a] = \lambda_a^{\eta} / \sum_v \lambda_{av}^{\eta}$ for each entity a

set $\mathbb{E}[\gamma_t] = \lambda_t^{\gamma} / \sum_v \lambda_{tv}^{\gamma}$ for each time interval t

set $\mathbb{E}[\phi_{ak}] = \lambda_{ak}^{\phi,s} / \lambda_{ak}^{\phi,r}$ for each entity a and general topic k

set $\mathbb{E}[\xi_a] = \lambda_a^{\xi,s} / \lambda_a^{\xi,r}$ for each entity a

set $\mathbb{E}[\psi_t] = \lambda_t^{\psi,s} / \lambda_t^{\psi,r}$ for each time interval t

end

return λ

$f(t_d, t) \mathbb{E}[\epsilon_{dt}]$	Date	Author Entity	Subject
6.86	1976-07-07	Cairo	Possible SC meeting on Israeli rescue operation
6.18	1976-07-10	Kuwait	Media reaction to Bicentennial summary
6.15	1976-07-06	Damascus	Syria condemns Israeli operation to free Air France ...
5.91	1976-07-08	Tel Aviv	Passengers comment on Air France hijacking
5.89	1976-07-06	Stockholm	Possible SC meeting on Israeli rescue operation
5.38	1976-07-09	Nicosia	Bicentennial activities in Cyprus
5.09	1976-07-11	State	Security Council debate on Entebbe events CONFID...
4.77	1976-07-09	State	Travel of Peter M. Storm, House Budget Committee
4.76	1976-07-06	Jidda	Weekly Saudi Editorial Summary (June 30-July 6)
4.68	1976-07-08	Lusaka	SWAPO President seeks assessment of Kissinger-Vor...
4.56	1976-07-07	Stockholm	Ugandan role in Air France hijacking
4.45	1976-07-06	Karachi	Transitional quarter funding for RSS travel
4.43	1976-07-06	Athens	Bicentennial anniversary in Greece
4.37	1976-07-08	Damascus	Beirut travel
4.34	1976-07-10	State	Status of Mrs. Bloch
4.17	1976-07-07	Hong Kong	Hong Kong Communist press denounces Israeli resc...
4.12	1976-07-08	Dar es Salaam	President Nyerere's fourth of July messages
4.09	1976-07-10	Moscow	Pravda and Krasnaya Zvezda on Entebbe rescue oper...

Table 5: Highest-ranked messages for the week immediately following the U.S. Bicentennial Celebration and Operation Entebbe. Capsule accurately recovers messages related to both of these real-world events. Typos are intentionally copied from the data.

References

- Michael I. Jordan, Zoubin Ghahramani, Tommi S. Jaakkola, and Lawrence K. Saul. 1999. An introduction to variational methods for graphical models. *Machine Learning*, 37(2):183–233, November.
- Martin J. Wainwright and Michael I. Jordan. 2008. Graphical models, exponential families, and variational inference. *Foundations and Trends in Machine Learning*, 1(1-2):1–305, January.

Top Terms
church, vatican, catholic, bishop, pope, ford, cardinal, ban, religious, archbishop program, university, grant, education, school, post, institute, research, center, american security, council, terrorist, threat, sc, sabotage, protective, herein, unsc, honour visit, hotel, schedule, arrival, arrive, depart, please, meet, day, room labor, union, strike, ilo, employment, federation, afl cio, trade, worker, confederation bank, credit, loan, investment, finance, payment, financial, eximbank, opic, central law, case, court, legal, investigation, arrest, justice, sentence, trial, attorney party, government, election, opposition, national, leader, campaign, vote, support, anti tax, company, pay, lease, compensation, exemption, repatriation, income, taxation, fee oil, petroleum, opec, crude, gulf, price, exploration, refinery, energy, company israel, arab, israeli, middle, egypt, peace, plo, cairo, egyptian, lebanon radio, television, broadcast, allotment, appropriation, obligation, zero, warc, transmitter, network india, indian, pakistan, delhi, goi, ocean, bangladesh, transit, pakistani, afghan turkish, turkey, cyprus, greek, greece, athens, ankara, morocco, cypriot, algeria aid, relief, emergency, usaid, disaster, donor, wfp, sahel, ifad, unicef aircraft, team, flight, clearance, transport, civair, aviation, traffic, charter, cargo soviet, moscow, press, ussr, soviet union, american, one, war, communist, article sea, zone, marine, maritime, fish, coastal, continental, territorial, mile, fishery

Table 6: The highest-probability vocabulary terms for a selection of general topics (one per row) according to $\mathbb{E}[\beta_1], \dots, \mathbb{E}[\beta_K]$. These examples come from the analysis described in [section 5](#). Capsule identifies diplomatic themes that are relevant to any entity.

Entity	Top Terms
Ankara	turkish, turkey, ankara, government, cyprus, greek, party, one, time
Athens	greek, athens, greece, gog, government, cyprus, turkish, press, minister
Auckland	new zealand, company, box, trade, contact, opportunity, united states
Baghdad	iraqi, iraq, goi, arab, state, regime, ministry, government, party
Berlin	berlin, frg, german, senat, time, bonn, trade, one, agreement
Bern	swiss, bern, federal, bank, snb, gold, end, interest, national
Brussels	belgian, belgium, brussels, government, firestone, european, ministry
Budapest	hungarian, hungary, trade, mudd, one, time, puja, well, policy
Buenos Aires	argentine, argentina, goa, us, hill, government, one, press, police
Cairo	egyptian, cairo, egypt, arab, israeli, israel, peace, agreement, president
Canberra	australian, australia, goa, government, minister, whitlam, end, dfa, time
Dakar	senegalese, president, african, summary, conference, end, support, one
Dar es Salaam	tangov, salaam, tanzanian, spain, president, government, african, one
Guayaquil	ecuador, ecuadorean, port, congen, one, tuna, local, time, boat
Islamabad	pakistan, gop, government, one, party, minister, general, opposition, ppp
Paris	paris, france, rush, french, one, government, amconsul, quai, european
Jerusalem	jerusalem, bank, israeli, us, israel, plo, one, arab, unifil
Jidda	saudi, jidda, saudi arabia, prince, us, fahd, one, time, government
Johannesburg	black, africa, african, trade, union, police, labor, one, committee
Kabul	afghan, government, goa, minister, one, pakistan, regime, time, ministry
Lima	peru, gop, lima, peruvian, dean, minister, general, marcona, government
Lisbon	portugal, portuguese, gop, lisbon, government, party, summary, minister
London	london, british, government, fco, labor, agreement, one, washdc, summary
Madrid	spanish, spain, madrid, one, govt, general, committee, government, time
Nairobi	kenya, nairobi, marshall, embassy, kenyan, unep, le, ref, state
Oslo	norwegian, norway, soviet, government, minister, ministry, policy
Ottawa	canadian, canada, goc, ottawa, us, extaff, government, minister, federal
Peking	chinese, peking, uslo, china, people, teng, one, trade, delegation, hong
Phnom penh	penh, phnom, khmer, rice, fank, enemy, cambodia, government, dean
Prague	czechoslovak, goc, czech, trade, embassy, one, mfa, time, cssr
Quito	ecuador, ecuadorean, gulf, government, minister, bloomfield, general, one
Sao Paulo	paulo, brazil, state, brazilian, president, government, congen, one, do
Seoul	korea, korean, rok, rokg, seoul, park, government, president, time
Singapore	singapore, asean, minister, government, one, prime, comment, vietnam
Sofia	bulgarian, trade, one, agreement, american, visit, committee, party
Sydney	australia, australian, one, general, american, state, government, post
Tokyo	japan, japanese, tokyo, fonoff, summary, miti, end, diet, time
Taipei	taiwan, groc, china, chinese, government, american, one, local, republic
The Hague	dutch, netherlands, hague, government, minister, party, stoel, mfa, one
USUN New York	committee, usun, priority, report, draft, resolution, sc, comite, rep, new york
Vancouver	canada, government, canadian, british, columbia, pipeline, federal, editorial
Zagreb	yugoslav, yugoslavia, croatian, fair, belgrade, american, one, ina, summary
Zurich	swiss, congen, consulate, general, american, bern, dollar, shipment

Table 7: The highest-probability vocabulary terms for a selection of entity topics (one per row) according to $\mathbb{E}[\eta_1], \dots, \mathbb{E}[\eta_A]$. These examples come from the analysis described in [section 5](#). Capsule identifies themes and interests that are specific to the entities.

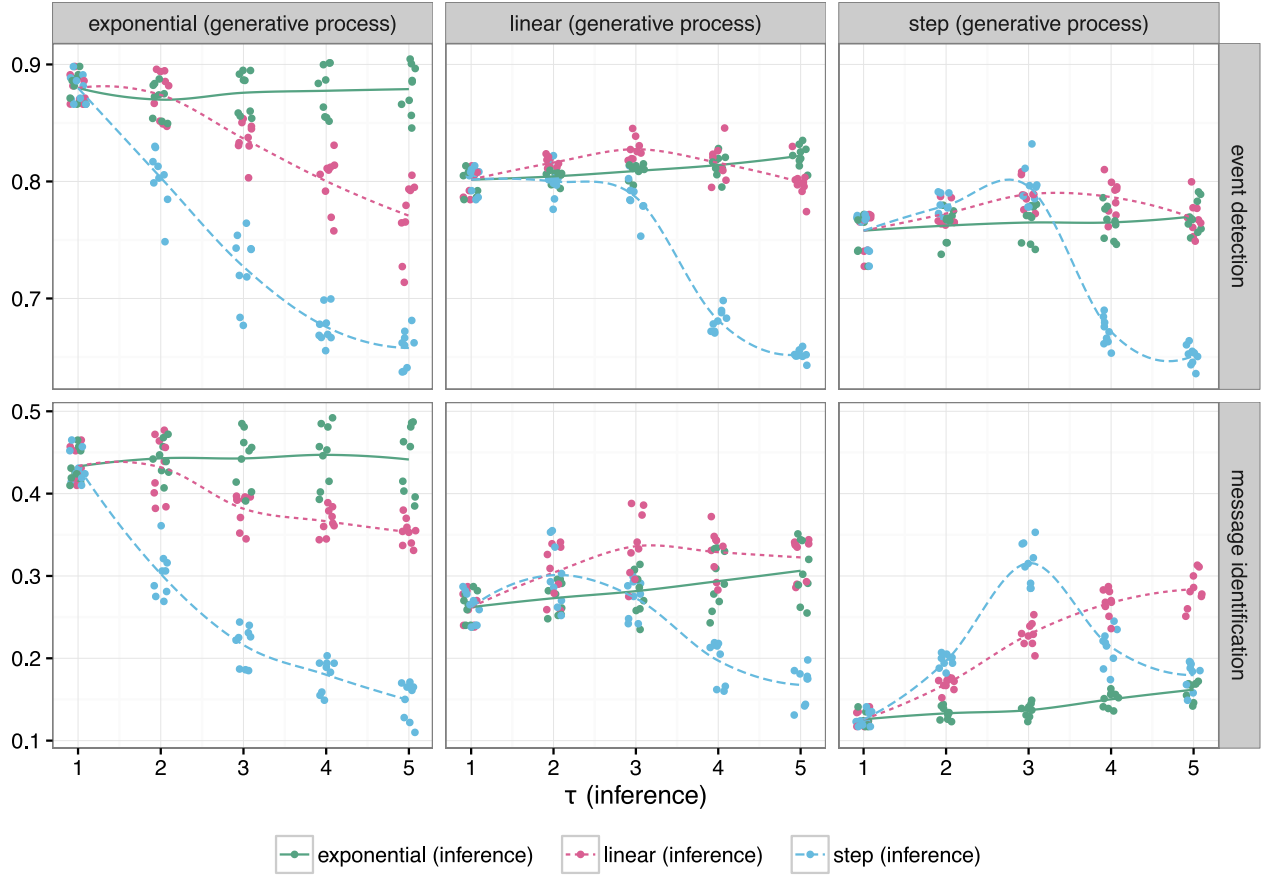


Figure 7: Capsule's sensitivity to several different decay functions (exponential, linear, and step) using simulated data. Capsule performs best when the decay function used in the inference algorithm matches the decay function used to generate the data. The exponential function is least sensitive to the value of τ used to generate the data and the value of τ used in the inference algorithm.